

Virtual Memory

- ▶ So far when discussing memory addresses we haven't been especially realistic
- ▶ When code is compiled, the addresses of variables are hardcoded into the executable
- ▶ If these were the actual physical locations in memory this would make running two copies of the code at the same time difficult (impossible!)
- ▶ A value updated in one process would modify the physical memory which was also being used by the second copy
- ▶ Obviously, since we can run multiple copies of the same executable concurrently this is not what happens

Virtual Memory

- ▶ Virtual memory adds a layer of abstraction between the process and the hardware
- ▶ We translate the addresses in your code (virtual addresses) into addresses locating information in real memory (physical addresses)
- ▶ This mapping is different for each executing process
- ▶ The Memory Management Unit (MMU) performs these translations in a transparent way for the executing program

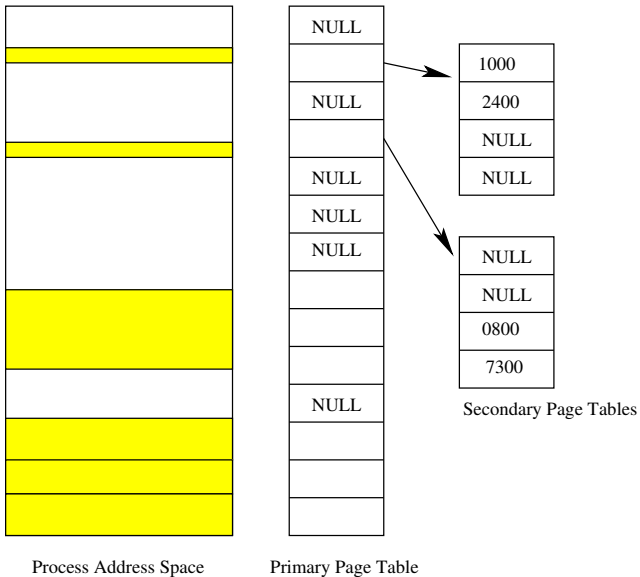
Memory Pages and Page Tables

- ▶ Physical memory is divided into fixed sized chunks called pages
- ▶ Size of pages depends on the hardware — typical sizes are 512 bytes to 4K
- ▶ All addressable locations in memory are in a page
- ▶ Therefore each location can be addressed by a page number and an offset in the page
- ▶ The page tables contain the mapping from virtual memory to physical memory
- ▶ Two separate page tables for kernel and user activities
- ▶ A set of page tables for each executing process

The Page Table

- ▶ Most modern page tables are arranged in a multi-level hierarchy
- ▶ This hierarchy saves spaces but has a small performance hit
- ▶ Useful because not every part of a processes address space needs to be mapped to physical memory
- ▶ Most accesses in memory are adjacent — on the heap, the stack (more on these later)
- ▶ In a 32-bit system you might have a primary table with 2^{12} entries each of which point to a secondary table with 2^8 each pointing at a 2^{12} byte page of physical memory
- ▶ The page table also stores some meta-data about the page including clean/dirty, present, process id etc.

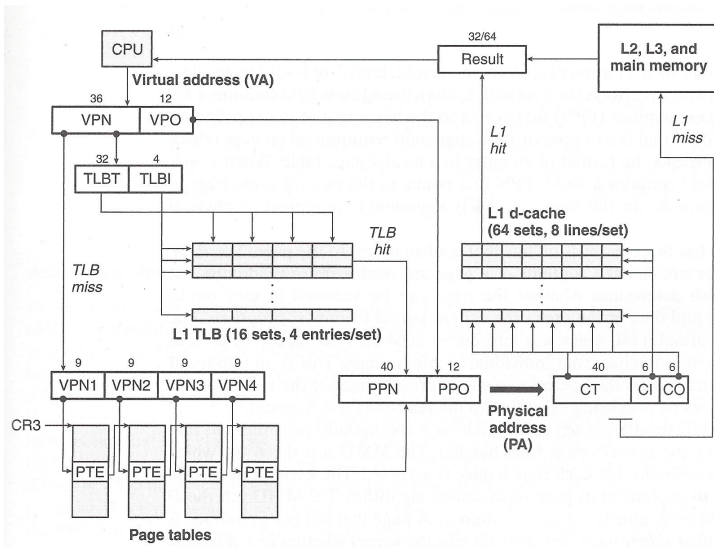
Page Tables



Translation Lookaside Buffer

- ▶ To reduce the cost of lookups the translation lookaside buffer (TLB) stores recent page table translations
- ▶ Implemented in hardware (see over)
- ▶ Usually achieve hit rates of 95%
- ▶ Linear memory accesses are nice to the TLB. Exploit locality of reference yet again!

Translation Lookaside Buffer




Processes

- ▶ A process is the execution of a program
- ▶ Consists of three types of data
 - ▶ *text* the CPU instructions
 - ▶ *data* where malloc'ed variables are kept
 - ▶ *stack* where local variables and other housekeeping is kept
 - ▶ Actually two stacks — user and kernel
- ▶ Operates in either user mode or kernel mode
- ▶ In kernel mode the process has full control of the system
- ▶ Can only get into kernel mode through well defined system calls


Executables

- ▶ Four regions
- ▶ Symbols — used for debugging
- ▶ Initialised data — global variables with values
- ▶ Text — the actual instructions to be executed
- ▶ Header — basic information needed to start the process




Processes

- ▶ Six regions
- ▶ Kernel Stack
- ▶ User Stack
- ▶ Heap
- ▶ Uninitialised data
- ▶ Initialised data
- ▶ Text



Stack Frames

- ▶ Each function that is called generates a new frame on the stack which contains
 - ▶ The return address from the function
 - ▶ Local variables
 - ▶ Parameters to the function
- ▶ Say we have a function `conj()` which calls `dot()`



The Stack

- ▶ The stack has a maximum size that is set as the program starts executing
- ▶ We can examine this using `ulimit -s`
- ▶ If we use more than this amount of memory our program will crash. This is called stack overflow
- ▶ Two main ways to do this
 - ▶ Declare large stack variables. Create a local array such as `double x[100000]`
 - ▶ Too many function calls. Each function call uses a small amount of memory for parameters, saved variables and a return value. Infinite recursion is usually to blame here.

Growing the heap

- ▶ When you allocate more memory using `malloc()` it appears on the heap
- ▶ The system uses `brk()` or `sbrk()` to extend the size of the heap
- ▶ `brk()` performs some checks and then calls `growreg()` to create new page tables for the region (if required)
- ▶ Passing `sbrk()` a negative argument decreases the size of the heap and releases page tables that are no longer required
- ▶ Can use `sbrk(0);` to find out where the heap currently ends

Page Faults

- ▶ Page faults occur when the MMU finds the page table entry for the required page is invalid. This isn't always an error
- ▶ May have to
 - ▶ Load a page from disk
 - ▶ Zero the existing page
 - ▶ Kill the process
- ▶ Since a page fault involves talking to the kernel and a possible disk transfer, they are expensive

Page Faults

- ▶ Several types of page faults can occur
- ▶ DISK — need to look on disk at a given location for the data
- ▶ NULL — physical memory not yet allocated
- ▶ MEM — protection level fault
- ▶ IOP — wait for I/O to complete
- ▶ LOCK — trying to modify a locked page

VM Benefits

- ▶ Protection — by only mapping pages belong to a process you can prevent access to another's data
- ▶ Demand paging — loading an entire program at once is wasteful. Demand paging only loads parts of a program as needed
- ▶ Shared pages — two processes running the same program can make use of the same pages of an executable
- ▶ Memory Mapped Files — a file can be made to look like part of a process' address space. When a page is read it is first copied in from the file. Used for shared libraries

Compilation Process

- ▶ Generating an executable consists of four phases
 - ▶ Pre-processing
 - ▶ Compilation
 - ▶ Assembling
 - ▶ Linking

Compilation Process

- ▶ The pre-processing phase involves all the hash statements in the program - `#include`, `#define`, `#pragma`
- ▶ The output from this phase can be seen using `gcc -E`
- ▶ Possible to do clever things here with `#ifdef` statements
- ▶ Pass arguments to the compiler using `gcc -DFOO`

Compilation Process

- ▶ Compilation takes the output from the pre-processor and turns it into assembly code
- ▶ This itself is a five stage process
 1. Lexical Analysis
 2. Parsing
 3. Semantic Analysis
 4. Optimization
 5. Code Generation
- ▶ Assembly code can be generated by using `gcc -S`

Compilation Process

- ▶ Assembling turns the assembly code into machine instructions specific for the processor architecture
- ▶ Normally just a straight translation
- ▶ There might be some further minor changes to the code at this stage (reordering, padding)
- ▶ Object code is generated using `gcc -c`

Compilation Process

- ▶ The final stage of generating an executable is linking
- ▶ This allows multiple object files to be combined into a single program
- ▶ Linking uses the `ld` command, but most of the time we just let `gcc` handle it for us
- ▶ Linking is also the time when references to external libraries are resolved
- ▶ Two forms of linking
 - ▶ Static linking - all the object code and libraries are rolled into a single executable
 - ▶ Dynamic linking - the object code is merged but references to libraries are left as pointers

Shared Libraries

- ▶ Shared libraries are loaded at the start of execution of a program
- ▶ They contain the program code that is to be added to the executable
- ▶ Shared libraries live in pre-defined locations on the filesystem — check `/etc/ld.so.conf`
- ▶ In modern systems this file usually includes additional files from a sub-directory
- ▶ Can add additional locations by modifying the `LD_LIBRARY_PATH` environment variable

Shared Libraries

- ▶ Shared libraries have three 'names'
 - ▶ real name: libgmp.so.3.5.0
 - ▶ soname: libgmp.so.3
 - ▶ compiler name: gmp
- ▶ The soname is used by the linker to determine which version of a library to use
- ▶ The compiler name is what you use as part of your `gcc` command
- ▶ `gcc -o foo foo.c -lgmp`

Shared Libraries

- ▶ To make your own shared libraries you recompile your code to generate it as position independent (or relocatable) code.
- ▶ Use either `gcc -fpic` or `gcc -fPIC`
- ▶ The position of the code in your process is defined in the global offset table (GOT)
- ▶ `-fpic` uses a machine-specific maximum size for the GOT. If your library is large your code might exceed the GOT
- ▶ `-fPIC` ignores the maximum GOT size but ends up producing larger code than `-fpic`

Compilation Process

- ▶ Once an executable has been created you can examine what shared libraries it requires by running `ldd a.out`

```
[dfrost@chuck ~]$ ldd ./a.out
linux-vdso.so.1 => (0x00007fffe87c9000)
libc.so.6 => /lib64/libc.so.6 (0x0000003fe1600000)
/lib64/ld-linux-x86-64.so.2 (0x0000003fe1200000)
```

- ▶ Can replace these references at runtime using the `LD_PRELOAD` mechanism
- ▶ Allows a user to replace function calls from libraries with their own version