

Interprocess Communication

- ▶ Single processes have limits on
 - ▶ Memory usage
 - ▶ Execution time
 - ▶ Disk quotas
- ▶ Processes can be executing on different machines that want to share information
- ▶ Need to be able to pass data between different processes
- ▶ Have already used some high level, complex forms of IPC, namely MPI, pthreads and OpenMP
- ▶ Want to look at more primitive types

Signals

- ▶ Signals are a primitive method to tell processes when an asynchronous event occurs
- ▶ Processes can send each other signals using the *kill(...)* system call
- ▶ You can only send signals to processes that are running with the same user id (unless you are root)
- ▶ Other signals are sent by the operating system when errors or events occur
 - ▶ Invalid memory access
 - ▶ Lack of system resources
 - ▶ Terminal interaction
 - ▶ Tracing process execution

Standard Signals

- ▶ HUP — hang up from a controlling terminal
- ▶ INT — interrupt from keyboard
- ▶ QUIT — quit (CTRL-C)
- ▶ TRAP — breakpoint reached
- ▶ BUS — bad memory access
- ▶ FPE — floating point exception
- ▶ ILL — illegal instruction
- ▶ KILL — terminate process
- ▶ SEGV — invalid memory reference
- ▶ STOP — suspend execution
- ▶ CONT — continue execution
- ▶ TSTP — stop from the keyboard (CTRL-Z)

Handling signals

- ▶ Each process has its own signal table which tells it how to deal with signals
- ▶ The default is to call *exit()*
- ▶ Using the *signal(...)* system call a user can designate another function to be executed instead of *exit()*
- ▶ The function in the table is executed immediately on receipt of the signal (once the program leaves kernel mode)
- ▶ Once the signal has been handled (assuming *exit* hasn't been called) the process continues executing as if nothing happened

Signal Uses

- ▶ Used by the scheduler to organise process execution
- ▶ Used by debuggers to control the process execution
- ▶ Print out current process status or important data structures
- ▶ Programs such as web servers, mail servers etc. often use the HUP signal to force a re-read of the configuration files

File descriptors

- ▶ Usually when we want to open a file for reading or writing we use `fopen(...)` and then `fprintf(...)` and `fscanf(...)` or `fgets(...)`
- ▶ This gives us a pointer to a `FILE` structure
- ▶ However the system calls for I/O use a file descriptor rather than a file pointer
 - ▶ `int open(const char *pathname, int flags);`
 - ▶ `int close(int fd);`
 - ▶ `ssize_t read(int fd, void *buf, size_t count);`
 - ▶ `ssize_t write(int fd, const void *buf, size_t count);`
- ▶ In the UNIX world, everything gets treated like a file — files, network access, directories, disks
- ▶ Use the same low level system calls to access them

File descriptor table

- ▶ Each process has a table which provides a mapping between numbers and open files
- ▶ The file descriptor is the id in the table
- ▶ By default every process has three open file descriptors at startup
 - ▶ 0 — stdin
 - ▶ 1 — stdout
 - ▶ 2 — stderr
- ▶ These are the file descriptor numbers

Pipes

- ▶ The next simplest form of IPC are pipes
- ▶ Found in all flavours of Unix (and other OSes)
- ▶ They allow data transfer in a FIFO manner
- ▶ Two kinds of pipes
 - ▶ Named pipes
 - ▶ Unnamed pipes
- ▶ Named pipes are created using `mknod`
- ▶ Unnamed pipes are created using `pipe`

Named pipes

- ▶ Named pipes are accessed like files
- ▶ Use `open`, `read`, `write`, `close`
- ▶ They are accessible to all processes with relevant permissions to access it
- ▶ The problem is that you need to know where in the filesystem the named pipe is

```
% mknod mypipe p  
% cat /etc/passwd > mypipe
```

```
% cat < mypipe  
root:x:0:0:root:/root:/bin/bash  
daemon:x:1:1:daemon:/usr/sbin:/bin  
:  
:
```

Unnamed pipes

- ▶ Unnamed pipes are only accessible by related processes
- ▶ Normally a process calls `pipe` and then uses `fork` to create a child to which it wants to send data
- ▶ `pipe(int fd[2])` creates a pair of file descriptors
 - ▶ `fd[0]` is used for reading
 - ▶ `fd[1]` is used for writing

```
int fd[2];
pipe(fd);
pid = fork();
if(pid)
    write(fd[1], string, strlen(string));
else
    read(fd[0], buffer, buflen);
```

Unnamed pipes

- ▶ Unnamed pipes occur relatively frequently
- ▶ `ls -al | sort -n +4 | less`
- ▶ Almost all unix pipes are uni-directional
- ▶ To get two-way communication you need to make two pipes
- ▶ Normally you close the unused ends of the pipes in the relevant process
- ▶ Using `dup` with pipes can lead to a nice interaction with `stdin` and `stdout`

Shared Memory

- ▶ Comes from SystemV Unix
- ▶ Processes can mark regions of their heap as shareable
- ▶ Another process can attach to and use this region as part of their own address space
- ▶ Once attached successfully, the second process uses the memory as normal — no further system calls are required
- ▶ Each shared memory segment has a key associated with it
- ▶ These are unique across the system

Shared Memory

- ▶ Creating a shared region
- ▶ `shm = shmget(32, 128*1024, 0777|IPC_CREATE);`
- ▶ Attaching a shared region to the address space of another (or the same) process
- ▶ `addr = shmat(shm, 0, 0);`
- ▶ Detaching the region
- ▶ `shmdt(addr);`
- ▶ Controlling the region
- ▶ `shmctl(shm, IPC_RMID, 0);`

Shared Memory

- ▶ The key to be used has to be shared between the processes somehow — chicken and egg problem
- ▶ The kernel cannot know which processes are allowed access to a shared memory segment
- ▶ The sharer can only grant read/write access as in file access
- ▶ Rogue processes can just guess the right ids and corrupt data

Sockets

- ▶ As outlined, pipes and shared memory both have drawbacks and limitations
- ▶ To overcome some of these we introduce sockets
- ▶ First appeared in BSD Unix
- ▶ Goals
 - ▶ Transparency — communication shouldn't depend on whether the processes are on the same machine or not
 - ▶ Efficiency — avoid needing server processes to regulate the communication
- ▶ Sockets rely on the TCP/IP protocols for communication

OSI Model

- ▶ An international standard that separates the different levels of abstraction into layers
- ▶ Each layer performs a well defined function
- ▶ The model has seven layers
 - ▶ Application — SMTP, HTTP etc.
 - ▶ Presentation — encoding of data
 - ▶ Session — manage dialogue control
 - ▶ Transport — ensures pieces are presented in order
 - ▶ Network — routing across the network
 - ▶ Data link — error free local transmission
 - ▶ Physical — electronics of high and low voltage
- ▶ This is a reference model — nobody actually implements this to manage their communication

TCP/IP basics

- ▶ The TCP/IP suite of protocols implements the network and transport layers of the OSI model
- ▶ There are two modes of operation
 - ▶ connection oriented, reliable (TCP)
 - ▶ connectionless, unreliable (UDP)
- ▶ TCP (Transmission control protocol) breaks up your message into small chunks for transfer over the network and re-assembles them in the right order at the other end
- ▶ UDP (User datagram protocol) allows small messages to be delivered quickly but with no checking for out of order, missing parts, duplicate arrivals etc.

Network addresses

- ▶ To allow programs to talk to each other, each application and machine must be individually addressable
- ▶ Each address consists of two parts
 - ▶ IP address to identify the machine
 - ▶ TCP/UDP port to identify the application
- ▶ IPv4 addresses are of the form 134.226.113.65. This gives 2^{32} addresses
- ▶ TCP/UDP ports are a number between 1 and 65536
- ▶ A new version of the IP protocol, IPv6, uses 128 bits and so gives 2^{128} addresses - approx 1×10^{28} addresses per person on the planet
- ▶ The DNS service provides a mapping between IP addresses and human memorable host names (woodstock.tchpc.tcd.ie)

Sockets

- ▶ The normal method for network programming is to use BSD sockets
- ▶ The end points (IP address and port) are represented by a socket
- ▶ The *socket(...)* system call creates a socket which is a file descriptor. It is used by all the rest of the networking calls
- ▶ Next we need to choose which port we are using
- ▶ If the program is the server then we need to explicitly bind to a specific port using the *bind(...)* call
- ▶ If the program is a client then we normally don't care which port we use and leave it to the OS to pick one

Sockets

Server

- ▶ Use *listen(...)* to indicate that new connections are to be accepted by the socket
- ▶ Connections are received one at a time using *accept(...)*

Client

- ▶ Use *connect(...)* to establish communication to the server process

Sockets

- ▶ Once a connection has been established processes may use *read/write* or more commonly the *send/recv* functions
- ▶ Sockets are bi-directional — no need to open two sockets like we needed with pipes
- ▶ A socket is just a file descriptor. We can use *fdopen(...)* to get a file handle
- ▶ Can then use *fprintf(...)* and *fscanf(...)* to send and receive messages