

MPI Status object

- ▶ We have been putting an MPI_Status object at the end of our Recv calls
- ▶ This object contains three elements
 - ▶ MPI_SOURCE — the sender of the message
 - ▶ MPI_TAG — the tag the sender used
 - ▶ MPI_ERROR — any error code from the Recv
- ▶ We can also determine the number of items received using a function call
- ▶ MPI_Get_count(*stat, type, *count)

MPI Status object

- ▶ We use the Status object when our receive uses wildcards such as `MPI_ANY_TAG` or `MPI_ANY_SOURCE`
- ▶ We may still want to know who sent us some data in order to communicate back to them
- ▶ In some communications the amount of data being received may not be known a priori
- ▶ Will need to know how much of our array contains valid data

Collective Operations

- ▶ Collective communications occur on all processes in the communicator
- ▶ Since we are just using `MPI_COMM_WORLD` that means all processes in the calculation
- ▶ Three main classes
 - ▶ Global synchronization
 - ▶ Global communications
 - ▶ Global reductions

Collective Operations

- ▶ Collectives insist that quantity of data being sent matches that being received
- ▶ Collectives are blocking - if one MPI task doesn't call the function then everyone else waits
- ▶ There are no tags so all calls to collective operations are matched in strict calling order

Collective Synchronization

- ▶ We have already met this
- ▶ `MPI_Barrier(comm);`
- ▶ All tasks in the communicator wait until everyone has called the barrier function

Collective Communication

- ▶ Three main forms
 - ▶ Root process sends data to everyone — broadcast and scatter
 - ▶ Root receives data from everyone — gather
 - ▶ Each process communicates with everyone else — allgather and alltoall

MPI_Bcast


- ▶ `MPI_Bcast(start, count, datatype, root, communicator);`
- ▶ The process with rank *root* sends the data to all other processes in the communicator
- ▶ These processes store the data at start
- ▶ After this operation each node has a full copy of all the data

MPI_Scatter

- ▶ `MPI_Scatter(send, sendcount, sendtype, recv, recvcount, recvtype, root, communicator);`
- ▶ The root process sends different sections of the send buffer to each process
- ▶ All processes (including the root) put this data in the recv buffer
- ▶ `sendcount` and `recvcount` are usually identical as are `sendtype` and `recvtype`
- ▶ MPI insists that $sendcount \times sizeof(sendtype) = recvcount \times sizeof(recvtype)$

MPI_Scatter

- ▶ Example — on root process we have *int send[100]*, on each proc *int recv[25]* and our simulation has 4 proc
- ▶ `MPI_Scatter(send, 25, MPI_INT, recv, 25, MPI_INT, 0, MPI_COMM_WORLD);`

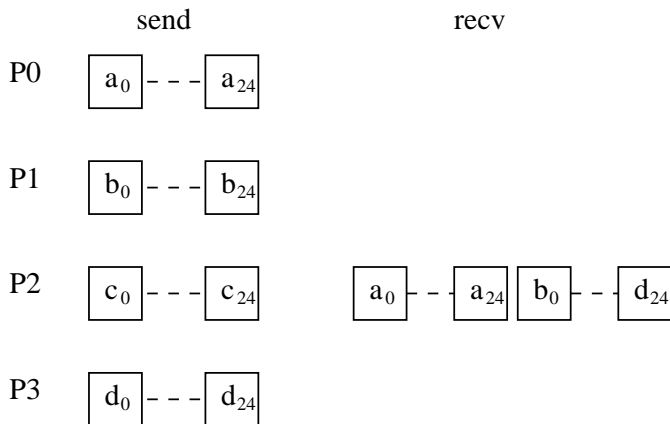


MPI_Gather

- ▶ MPI_Gather does the opposite of MPI_Scatter
- ▶ MPI_Gather(send, sendcount, sendtype, recv, recvcount, recvtype, root, communicator);
- ▶ The root process receives copies of the send buffer from each process
- ▶ The data is assembled in rank order in the recv buffer

MPI_Gather

- ▶ `MPI_Gather(send, 25, MPI_INT, recv, 25, MPI_INT, 2, MPI_COMM_WORLD);`



MPI_Reduce

- ▶ `MPI_Reduce(send, recv, count, datatype, operation, root, communicator);`
- ▶ The data in the send buffer of each process is reduced by the operation
- ▶ These operations can be addition, multiplication, maximum, minimum etc.
- ▶ The result is stored in the recv buffer on the root process

All collectives

- ▶ MPI also provides versions of Reduce and Gather where the results happen on all processes
- ▶ `MPI_Allreduce(send, recv, count, type, operation, communicator);`
- ▶ `MPI_Allgather(send, sendcount, sendtype, recv, recvcount, recvttype, communicator);`
- ▶ Note there are no root processes in these calls
- ▶ In the *non all* versions the recv buffers on the non root processes can be NULL
- ▶ In the *all* versions each process must have sufficient space in the recv buffer

MPI_Alltoall

- ▶ Finally, there is a function that allows all processes to send different messages to every process at the same time
- ▶ `MPI_Alltoall(send, sendcount, sendtype, recv, recvcount, recvtype, communicator);`
- ▶ Each process sends `sendcount` items, starting from $rank \times sendcount$, from its send buffer
- ▶ Each process receives `sendcount` items from each other process and stores it in the recv buffer starting from $rank \times sendcount$
- ▶ Errors occur if the send and/or recv buffers are not large enough