

Amdahl's Law

- ▶ Gene Amdahl — designer of the IBM 360 mainframe
- ▶ Any task can be split into parts that can be parallelized and parts that can't
- ▶ Serial parts — setup of problem by reading in data, generating statistics after each iteration
- ▶ Parallel parts — numerical solver, integration of Newton's laws


Amdahl's Law

- ▶ Say we have task of which 95% can be executed in parallel
- ▶ Even if we use an infinite number of processes on the parallel part we still need 5% of the original time to execute the serial part
- ▶ Maximum speed up is given by the formula

$$\frac{1}{S + \frac{1-S}{N}}$$

where S is the proportion of the code to be executed in serial and N is the number of processes in the parallel part

Amdahl's Law



Amdahl's Law

- ▶ This suggests that there is no point in writing code for more than 8-16 processes
- ▶ Not true!
- ▶ As you run on larger problem sizes often the serial part scales linearly but the parallel part scales with n^2 or greater
- ▶ By tackling larger problems a 95% parallel problem can become a 99% parallel problem and eventually a 99.999% parallel problem
- ▶ Plugging in the figures a 99.9% problem on a 1024 way system gives a 506 speedup

Amdahl's Law

- ▶ Of course Amdahl's law assumes the parallel part is perfectly parallelizable
- ▶ It does not take into consideration the time spent passing data between processes
- ▶ There is a more complex form that takes transfer times and network latencies into account
- ▶ For our purposes the simple version is sufficient


Master-Slave model

- ▶ The Master-Slave model is one of the standard models in parallel programming
- ▶ It is used when the problem is *embarassingly parallel* i.e. each task is independant of the others
- ▶ Effectively a 100% parallel problem from Amdahl's perspective
 - ▶ Brute force attacks in cryptography
 - ▶ Non-realtime rendering for animation
 - ▶ Ray tracing
 - ▶ Pretty much all of bio-informatics
 - ▶ Genetic algorithms
- ▶ The model used in things like BOINC / *-at-home
- ▶ Sometimes called the manager-worker model to be slightly more politically correct

Master-Slave model

- ▶ One process in the calculation (usually rank 0) is designated the master
- ▶ It generates work and then allocates it to the other processes
- ▶ Each slave carries out its given task and sends the results back to the master
- ▶ If there is further work to do the master sends the slave a new task
- ▶ If there is no more work then the slave waits

Master-Slave model



Master-Slave model

- ▶ There is no mapping of work to processes — any task can be carried out by any process
- ▶ Need to make sure that the master is not a bottleneck. Can occur if
 - ▶ If the tasks are too small
 - ▶ Too many workers
- ▶ Choose size of tasks so the time taken to do the task is much larger than time taken to transfer work to slave and result back to master
- ▶ Waiting times are also reduced if the time taken for each task is somewhat random

Master-Slave model

Master

```
if(rank == 0) {  
    for(i=1;i<size;i++)  
        send_next_job(i);  
    while(jobs_left) {  
        get_result(s);  
        send_next_job(s);  
    }  
    for(i=1;i<size;i++) {  
        get_results(s);  
        no_more_jobs(s);  
    }  
}
```

Slave

```
} else {  
    while(1) {  
        get_next_job();  
        if(no_job)  
            break;  
        do_work();  
        send_result();  
    }  
}
```

Master-Slave model

- ▶ Need to have some agreed message to indicate the work has all been completed
 - ▶ Can use the tag field for this
- ▶ Master needs to know who has send him work
 - ▶ Use *MPI_ANY_SOURCE* in the recv
 - ▶ Use the *MPI_STATUS* object to determine the actual source

Genetic Algorithms

- ▶ Genetic algorithms are used to search a solution space for the optimum
- ▶ They are based on Darwinian "survival of the fittest" concepts
- ▶ The solution space is modelled using a string — the chromosome
- ▶ A population of chromosomes is generated
- ▶ They are selected according to their "fitness"
- ▶ They may crossover to form new offspring
- ▶ Random mutation may occur in new offspring

Genetic Algorithms

- ▶ Chromosomes are usually modelled as strings of 1s and 0s
- ▶ In C we usually use the bit representation of a char or int
- ▶ For initialisation we often just randomly fill the population
- ▶ Need to make sure they represent valid solutions to the problem
- ▶ The crossover operation is the key to GA
 - ▶ Choose a location in the string
 - ▶ Exchange the subsequences after this locus between two chromosomes to create two offspring
 - ▶ Given parents 10000100 and 11111111 to be crossed after the third bit
 - ▶ Offspring are 10011111 and 11100100

Genetic Algorithms

- ▶ Mutation occurs very rarely (0.1%) but are important in escaping local maximums
- ▶ The rate is sometimes varied as the population becomes more homogeneous
- ▶ Each iteration is called a generation
- ▶ Normally a GA is run for 100 or more generations

Genetic Algorithms

- ▶ Using chromosome of size 8 with a fitness function that counts the number of 1s in the string

Label	String	Fitness
<i>A</i>	00000110	2
<i>B</i>	11101110	6
<i>C</i>	00100000	1
<i>D</i>	00110100	3

- ▶ Might select B and D to be 1st set of parents and B and C to be the 2nd set — parents are chosen randomly by fitness (roulette wheel method)
- ▶ B and D crossover after first bit to give new offspring E (10110100) and F (01101110)
- ▶ B and C do not crossover

Genetic Algorithms

- ▶ Next mutation occurs. Say E is mutated at location 6 and B is mutated at location 1
- ▶ This gives us a population for the next generation

Label	String	Fitness
\hat{E}	10110000	3
F	01101110	5
C	00100000	1
\hat{B}	01101110	5

- ▶ Even though the highest fitness value is now 5 compared to 6 previously, the overall fitness of the population has increased from 12 to 14
- ▶ Eventually, if iterated often enough, this will result in a population of strings with all 1s

Genetic Algorithms

- ▶ Where are genetic algorithms useful?
- ▶ Quite often used in scheduling and timetabling problems
- ▶ Problems with a huge solution space — can't evaluate all possible solutions
- ▶ Problems with complex solutions spaces where normal techniques (simplex/hill climbing) might get stuck
- ▶ Protein folding and docking
- ▶ Travelling salesman problem

Prisoners' Dilemma

- ▶ A simple two-person game invented in the 1950s
- ▶ Two people (Alice and Bob) are held in separate cells in the cop shop accused of committing a crime
- ▶ A deal is offered to each — testify against the other, you will get off free and the other will get 5 years
- ▶ If both people take the deal then both have discredited the other's testimony and both get 4 years
- ▶ If both people don't take the deal then both are convicted of lesser charges and get 2 years each
- ▶ What should each person do?

Prisoners' Dilemma

- ▶ Create a reward matrix. Base this on the number of years off their prison time

	Co-operate	Defect
Co-operate	3,3	0,5
Defect	5,0	1,1

- ▶ What is the best strategy to maximize your own payoff?
- ▶ Take the deal! Your average time off is $\frac{5+1}{2} = 3$ as opposed to $\frac{3+0}{2} = 1.5$
- ▶ The game becomes more interesting if this process is iterated — you play several games in a row
- ▶ Both players always taking the deal leads to a much lower payoff than both not taking the deal
- ▶ How can you make your opponent co-operate for mutual benefit?
- ▶ Act based on how the previous games went

Prisoners' Dilemma

- ▶ Use a genetic algorithm to evolve strategies
- ▶ How to encode the problem into a string of bits?
- ▶ 4 possibilities from previous game
 - ▶ CC — both players co-operate with each other
 - ▶ CD — player 1 co-operates but player 2 stabs
 - ▶ DC — player 2 co-operates but player 1 stabs
 - ▶ DD — both players stab each other
- ▶ Invent a rule to apply based on what happened in the previous game
- ▶ If you look over the previous two games there are 16 possibilities
 - ▶ CC CC
 - ▶ CC CD
 - ▶ ...
 - ▶ DD DD

Prisoners' Dilemma

- ▶ To evaluate the fitness of each strategy, play a round-robin tournament between each strategy
- ▶ The fitness is the total payoff from all the games
- ▶ Then use this fitness value to select the next generation, perform cross-over and mutation and keep going

Parallel GA

- ▶ Genetic algorithms are very good targets for parallelising using the Master-Slave paradigm
- ▶ Generally fitness evaluation is the most expensive part of the GA
- ▶ Most of the time the evaluation of fitness can be split into independent operations
 - ▶ Counting bits in a string — each chromosome can be evaluated separately
 - ▶ Prisoners Dilema' — fitness of chromosomes are inter-dependant but matches between sets of chromosomes can be evaluated separately
- ▶ While the crossover/mutation code can be a bit messy once you have written it once it can be reused and it executes quickly