

MPI Datatypes

- ▶ Recall MPI has its own datatypes to wrap around primitive types
 - ▶ MPI_CHAR
 - ▶ MPI_INT
 - ▶ MPI_DOUBLE
- ▶ MPI also allows the user to define their own datatypes
- ▶ Useful for packing data together into a single MPI_Send()
- ▶ Also allows sending of non-contiguous memory in a single communication

Type Maps

- ▶ Derived datatypes are made up of sets of primitive datatypes
- ▶ They form what is called a *type map*
- ▶ The map consists of the type and the displacement (or offset) in bytes from the first element
- ▶ Say we have a datatype consisting of a double, an int and finally a char then the map is
- ▶ `{double, 0}, {int, 8}, {char, 12}`

Extent and Size

- ▶ The extent of a datatype is the distance between subsequent elements of the same datatype
- ▶ This can be different to the size of the datatype as certain variable may need to be aligned on 4 or 8 byte boundaries
- ▶ In the previous example the size of the data type is 13 bytes but the extent is 16 bytes
- ▶ This is because the double in a subsequent element could not start immediately beside the previous char due to memory alignment issues
- ▶ Use `MPI_Type_extent(type, *extent)` to determine

Contiguous Data

- ▶ Simplest derived datatype
- ▶ Consists of a number of contiguous items of the same datatype
- ▶ `MPI_Type_contiguous(count, oldtype, *newtype);`
- ▶ After we create a new datatype we must use `MPI_Type_commit(newtype)` before we use it for sending data
- ▶ When we are done we can use `MPI_Type_free(newtype)` to delete the datatype

Vector Data

- ▶ Vectors are similar to contiguous, but allow for gaps between the objects
- ▶ `MPI_Type_vector(count, block, stride, oldtype, *newtype);`
- ▶ There is a `hvector` version where stride is measured in bytes rather than objects.
- ▶ Useful for sending columns of matrix

```
MPI_Type_vector(nrows, 1, ncols, MPI_Double, &coltype);  
MPI_Type_commit(coltype);  
MPI_Send(&g[0][ncols-1], 1, coltype, dest, tag, comm);
```

Indexed Data

- ▶ A more generalized and less structured vector
- ▶ Allows sequences of blocks of different lengths with different offsets
- ▶ `MPI_Type_indexed(count, *blocks, *offsets, oldtype, *newtype);`
- ▶ Again a hindex version with offsets in bytes
- ▶ Not commonly used in scientific computing

MPI Structs

- ▶ General purpose wrapper for structured data
- ▶ `MPI_Type_Struct(count, *blocks, *offsets, *oldtypes, *newtype);`
- ▶ Remember that pointers within the struct won't be dereferenced.
- ▶ Just the value of the pointer will be sent which will most likely be useless on the receiving end.

MPI Structs

```
struct foo {  
    double x; int l, m; char c;  
};
```

```
MPI_Type_extent(MPI_DOUBLE, &exD);  
MPI_Type_extent(MPI_INT, &exI);  
MPI_Datatype type[3] = {MPI_DOUBLE, MPI_INT, MPI_CHAR};  
int blocks[3] = {1, 2, 1};  
int offsets[3] = {0, exD, exD+2*exI};  
  
MPI_Type_struct(3, blocks, offsets, type, &newtype);  
MPI_Type_commit(&newtype).
```


Packing and Unpacking

- ▶ Sometimes the data is not well organized to use vectors or structs
- ▶ Want to minimize the number of communications, so bundle data together into a single buffer before sending
- ▶ `MPI_Pack(input, count, type, output, size, position, comm);`
- ▶ Position is updated so subsequent calls to `MPI_Pack` will append data at the appropriate place
- ▶ Then use a single `MPI_Send` using `MPI_PACKED` datatype
- ▶ `MPI_Unpack(input, count, type, output, size, position, comm)` used to unpack the buffer on the recv'ing side

Summary

- ▶ User defined data types are great for sending non-contiguous data or data with a mixture of types in a single operation
- ▶ Remember to commit the new datatype before use
- ▶ These operation are not collective but must be run on each node where the datatype is to be used