

# Collective Communication

- ▶ Collective communications occur on all processes in the communicator
- ▶ Since we are just using `MPI_COMM_WORLD` that means all processes in the calculation
- ▶ Note - all processes must call the function otherwise other processes will become blocked

# Collective Communication

- ▶ `MPI_Bcast(...)` sends information to all processes
- ▶ `MPI_Reduce(...)` combines data from all processes and returns it to one process
- ▶ `MPI_Scatter(...)` splits up a large data set across all processes
- ▶ `MPI_Gather(...)` does the opposite of `MPI_Scatter`
- ▶ In some senses `MPI_Barrier()` is also a collective function even though no data is passed around

# MPI\_Bcast

- ▶ `MPI_Bcast(start, count, datatype, root, communicator);`
- ▶ The process with rank *root* sends the data to all other processes in the communicator
- ▶ These processes store the data at start
- ▶ After this operation each node has a full copy of all the data

# MPI\_Reduce


- ▶ `MPI_Reduce(send, recv, count, datatype, operation, root, communicator);`
- ▶ The data in the send buffer of each process is reduced by the operation
- ▶ These operations can be addition, multiplication, maximum, minimum etc.
- ▶ The result is stored in the recv buffer on the root process

# MPI\_Scatter

- ▶ `MPI_Scatter(send, sendcount, sendtype, recv, recvcount, recvtype, root, communicator);`
- ▶ The root process sends different sections of the send buffer to each process
- ▶ All processes (including the root) put this data in the recv buffer
- ▶ `sendcount` and `recvcount` are usually identical as are `sendtype` and `recvtype`
- ▶ MPI insists that  $sendcount \times sizeof(sendtype) = recvcount \times sizeof(recvtype)$

# MPI\_Scatter

- ▶ Example — on root process we have *int send[100]*, on each proc *int recv[25]* and our simulation has 4 proc
- ▶ `MPI_Scatter(send, 25, MPI_INT, recv, 25, MPI_INT, 0, MPI_COMM_WORLD);`

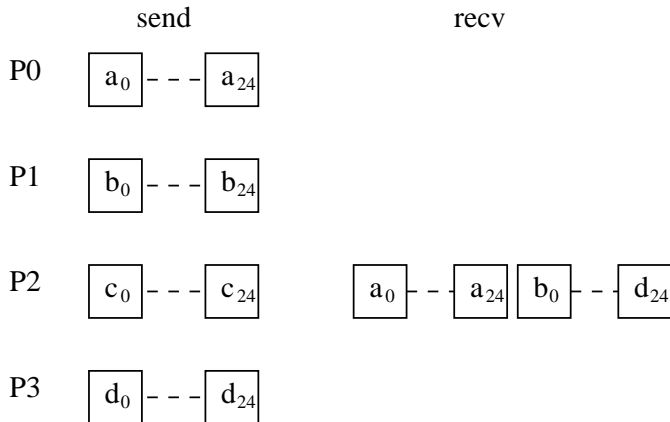


# MPI\_Gather

- ▶ MPI\_Gather does the opposite of MPI\_Scatter
- ▶ MPI\_Gather(send, sendcount, sendtype, recv, recvcount, recvtype, root, communicator);
- ▶ The root process receives copies of the send buffer from each process
- ▶ The data is assembled in rank order in the recv buffer

# MPI\_Gather

- ▶ `MPI_Gather(send, 25, MPI_INT, recv, 25, MPI_INT, 2, MPI_COMM_WORLD);`





# All collectives

- ▶ MPI also provides versions of Reduce and Gather where the results happen on all processes
- ▶ `MPI_Allreduce(send, recv, count, type, operation, communicator);`
- ▶ `MPI_Allgather(send, sendcount, sendtype, recv, recvcount, recvtype, communicator);`
- ▶ Note there are no root processes in these calls
- ▶ In the non all versions the recv buffers on the non root processes can be NULL
- ▶ In the all versions each process must have sufficient space in the recv buffer

# MPI\_Alltoall

- ▶ Finally, there is a function that allows all processes to send different messages to every process at the same time
- ▶ `MPI_Alltoall(send, sendcount, sendtype, recv, recvcnt, recvttype, communicator);`
- ▶ Each process sends `sendcount` items, starting from  $rank \times sendcount$ , from its send buffer
- ▶ Each process receives `sendcount` items from each other process and stores it in the `recv` buffer starting from  $rank \times sendcount$
- ▶ Errors occur if the `send` and/or `recv` buffers are not large enough

# MPI\_Sendrecv

- ▶ When passing data between processes in the Jacobi example we had to be careful to get the MPI\_Send and MPI\_Recv calls in the right order

Proc 0	Proc 1
MPI_Send(...)	MPI_Recv(...)
MPI_Recv(...)	MPI_Send(...)

- ▶ This type of communication pattern is very common
- ▶ MPI provides the MPI\_Sendrecv() function to get around the problem of deadlocking calls
- ▶ MPI\_Sendrecv(send, sendcount, sendtype, dest, sendtag, recv, recvcount, recvtype, source, recvtag, communicator, stat);
- ▶ Basically mash the send and receive functions into one call