


Sending modes in MPI

- ▶ We have claimed our MPI functions so far have been blocking
- ▶ That is they wait until the operation has been completed before exiting and allowing the program to continue



- ▶ Not exactly true. There are several modes for executing Sends and Recvs depending on buffering and synchronization wishes.

Sending modes in MPI

Synchronous MPI_Ssend Returns only when the matching receive occurs.

Buffered MPI_Bsend Returns immediately whether or not the transfer has occurred.

Ready MPI_Rsend Give an error if called before the matching receive occurs.

Standard MPI_Send Returns when you can safely use the original data.

All are matched by the same MPI_Recv call. No modes for Recv.

Synchronous Send

- ▶ `MPI_Ssend(...)`
- ▶ Full synchronization but potentially high overhead.
- ▶ Send may start whenever it wants.
- ▶ Recv must start before Send returns.
- ▶ Safest form of Sending. Maybe useful for porting code.

Buffered Send

- ▶ `MPI_Bsend(...)`
- ▶ Needs explicit buffer management by the user.
- ▶ `MPI_Buffer_attach(void *buf, int size)`
- ▶ Only one buffer allowed per MPI task at a time.
- ▶ Need to ensure buffer is big enough for all the data.
- ▶ `MPI_Buffer_detach(void *buf, int)` to detach the buffer.

Ready Send

- ▶ `MPI_Rsend(...)`
- ▶ Recieve must have started before Rsend is called.
- ▶ Useful on some specialized networks that can support handshakeless communications.
- ▶ Normally do not use this function!


Blocking Message Passing

- ▶ Unless the `MPI_Send(...)` and `MPI_Recv(...)` occur around the same time one of the processes will end up idling
- ▶ As mentioned previously these functions can also end up deadlocking if the ordering of the functions on each process is wrong — two `MPI_Send(...)` calls posted with no `MPI_Recv(...)`
- ▶ `MPI_Sendrecv(...)` can solve some of these problems but idling will still occur
- ▶ Solution is to use non-blocking operations

Non Blocking Message Passing

- ▶ In non blocking MPI, the functions return immediately after being called
- ▶ The data may not yet have been transferred
- ▶ Process is free to continue computation that does not depend on the data involved in the transfer
- ▶ When the data is required, perform a check to see if the non blocking call has been completed
- ▶ If so then free to continue computation
- ▶ Otherwise wait until the transfer is complete

Non Blocking Message Passing



Non Blocking Message Passing

- ▶ `MPI_Isend(data, count, type, dest, tag, comm, request);`
- ▶ `MPI_Irecv(data, count, type, src, tag, comm, request);`
- ▶ New structure `MPI_Request` contains the information about the state of the transfer
- ▶ Use `MPI_Test(request, flag, status)` or `MPI_Wait(request, status)` to find out if the transfer is completed
- ▶ If feeling brave (foolhardy) use `MPI_Request_free(request)` to abandon the checking

Non Blocking Message Passing

- ▶ Using `MPI_Isend` and `MPI_Wait` directly one after the other is equivalent to just using `MPI_Send`
- ▶ Doesn't add much, if any, overhead
- ▶ Normally just place the `MPI_Wait` just before the next call the needs to update the data
- ▶ We can use `MPI_Test` to build an event driven programme. This is often used on the Recv'ing side

Non Blocking Message Passing

- ▶ Usual rules about src, dest and tag apply — can use `MPI_ANY_SOURCE` and `MPI_ANY_TAG` in `MPI_Irecv(...)`
- ▶ `MPI_Isend(...)` can be received by any form of `MPI_Recv`
- ▶ `MPI_Irecv(...)` can get data from any form of `MPI_Send`

Multiple messages

- ▶ Sometimes we want to set up a group of non-blocking communications
- ▶ Create arrays of request and status objects
- ▶ Three options for waiting on multiple communications
- ▶ `MPI_Waitall(count, requests, statuses)`
- ▶ `MPI_Waitany(count, requests, index, statuses)`
- ▶ `MPI_Waitsome(count, requests, compcount, indices, statuses)`
- ▶ Similar functions for Test

Probing

- ▶ MPI also provides functions for you to see if there is a communication heading your direction
- ▶ `MPI_Probe(src, tag, comm, stat)`
- ▶ Allows us to allocate appropriate sized recv buffers.
- ▶ More efficient than having a super-sized catch all buffer.
- ▶ A non-blocking version `MPI_Iprobe(src, tag, comm, flag, stat)`
- ▶ Mixed views on how useful Probe is.

Cancelling

- ▶ Sometimes a non-blocking call will never be matched.
- ▶ Good code should clean up after itself.
- ▶ Use `MPI_Cancel(req)` to start the process.
- ▶ Then have to complete the request using wait/test as usual
- ▶ Finally use `MPI_Test_cancelled(stat, flag)` to see if the cancel was successful.

Golden Rule of Non Blocking Comms

- ▶ After `MPI_Isend(...)` do not change the data in the message until `MPI_Test(...)` or `MPI_Wait(...)` have indicated that it is safe to do so
- ▶ After `MPI_Irecv(...)` do not use the data in the message until `MPI_Test(...)` or `MPI_Wait(...)` have indicated that it is safe to do so

Persistent Communication

- ▶ If the same Send/Recv pair is called repeatedly can set up a persistent communication to avoid repeated message setup and teardown
- ▶ Saves on creation and destruction of MPI objects in the library
- ▶ Often useful in halo exchanges in data decomposition
- ▶ Three step process
 1. Create request objects
 2. Send and receive the messages using MPI_Start
 3. Delete the objects

Persistent Communication

```
MPI_Request send_req, recv_req;
MPI_Status stat;

MPI_Send_init(sendbuf, count, type, dest, tag, comm, &send_req);
MPI_Recv_init(recvbuf, count, type, src, tag, comm, &recv_req);

for (i=0; i<N; i++) {
    /* Halo exchange */
    MPI_Start(&recv_req);
    MPI_Start(&send_req);

    MPI_Wait(&send_req, stat);
    MPI_Wait(&recv_req, stat);

    /* Update values */
    do_stuff();
}

MPI_Request_free(send_req);
MPI_Request_free(recv_req);
```