


Threads intro


- ▶ We have mainly looked at writing parallel programs where each task has its own private memory space
- ▶ This model works well on distributed memory systems (clusters)
- ▶ You need to explicitly send data from one process to another



Threads intro

- ▶ In a machine with multiple processors it can be beneficial to have several processors all working on the same data set
- ▶ We use lightweight processes or threads to achieve this aim
- ▶ All programs we have written so far have only a single thread of execution
- ▶ Processes can have multiple threads executing at the same time
- ▶ These threads share the same global memory (heap) but each has its own stack

Threads intro



- ▶ Each processor in the system can access any physical memory location
- ▶ Some memory may be further away (NUMA)
- ▶ Sample machines
 - ▶ SGI Ultraviolet
 - ▶ Any multi-socket Intel/AMD server

Benefits of threads

- ▶ Portability — easy to move code from a serial to a parallel machine
- ▶ Latency — we can hide I/O and memory latencies by having one thread do work while another is waiting
- ▶ Load balancing — by allocating each task to a separate thread the system can load balance. No need for the programmer to explicitly schedule interaction
- ▶ Ease of use — threads are generally easier to write code for than MPI
- ▶ Future proof — modern machines are shipping with greater numbers of CPU cores. Threads will play a major role
- ▶ Major downside — need a machine with multiple cores to see any benefit

Creating threads

- ▶ Need to include `pthread.h`
- ▶ Need to link with `-lpthread`
- ▶ Use `pthread_create(...)` to create a new thread of execution
- ▶ Four arguments
 - ▶ `pthread_t *handle` — an identifier for the thread
 - ▶ `const pthread_attr_t *att` — attributes of the thread. Can use `NULL`
 - ▶ `void * (*function)(void *)` — the “main” function of the thread
 - ▶ `void * arg` — a single argument to the function. Again `NULL` acceptable

Ending threads

- ▶ All threads are terminated when `main()` exits
- ▶ An individual thread terminates when you call `pthread_exit` within the thread
- ▶ A thread also ends if you reach the end of the thread's main function
- ▶ Threads can be terminated by other threads calling `pthread_cancel`

Waiting for threads

- ▶ Need a way to wait for the threads to finish properly
- ▶ Use `pthread_join(...)`
- ▶ Two arguments
 - ▶ `pthread_t *handle` — the identifier of the thread you are waiting for
 - ▶ `void **ptr` — the value passed to `pthread_exit`
- ▶ Note multiple threads can't join a single thread
- ▶ Threads can become `detached` and cannot be joined. This should generally be avoided

Other management calls

- ▶ `pthread_self` gets your own thread id
- ▶ `pthread_equal` compares two threads. Cannot just use normal `==` comparison as thread objects are not basic data types
- ▶ `pthread_once(ctl, func)` executes `func` exactly one time per process. Every other call to the function has no effect.
- ▶ Often used for an initialization function

Shared Variables

- ▶ The best thing about threads is that the communications just happens
- ▶ The worst thing is that you need to synchronise the threads so they don't mess up the shared data
- ▶ Consider the following code

```
if( localval > globalval)  
    globalval = localval;
```

- ▶ If globalval is 100 and you have two threads with localval set to 150 and 200
- ▶ If both execute this section of code simulataneously then you might end up with globalval set to either 150 or 200

Mutexes

- ▶ The main method of avoiding these *race conditions* is using mutexes (mutual exclusions)
- ▶ A mutex has two states — locked and unlocked
- ▶ Only one thread can have a lock at a time
- ▶ If you try to lock a mutex that is already locked by another thread your execution is blocked
- ▶ When the other thread unlocks the lock then you can get the lock and use the shared data

```
Get_Lock();  
if( localval > globalval)  
    globalval = localval;  
Release_Lock();
```

Mutexes

- ▶ The pthread library defines a new data type called `pthread_mutex_t` to handle mutex locks
- ▶ Use the function `pthread_mutex_lock(*lock)` to acquire the lock
- ▶ If the mutex is already locked you will wait until it becomes unlocked
- ▶ Need to be careful about circular deadlocks!
- ▶ Use the function `pthread_mutex_unlock(*lock)` to release the lock
- ▶ Before using a mutex it needs to be initialised
- ▶ Use `pthread_mutex_init(*lock, *attrib)` to initialise. Just use NULL for the attributes for the moment.

Conditional Locks

- ▶ Using two mutexes we can avoid most deadlocks
- ▶ We need to request them in different orders and use `pthread_mutex_trylock` to try to grab the locks

Thread 1	Thread 2
<pre>pthread_mutex_lock(lock1); pthread_mutex_lock(lock2); /* Do calculations */ pthread_mutex_unlock(lock2); pthread_mutex_unlock(lock1);</pre>	<pre>while(1) { pthread_mutex_lock(lock2); if(pthread_mutex_trylock(lock1) == 0) break; pthread_mutex_unlock(lock2); } /* Do calculations */ pthread_mutex_unlock(lock1); pthread_mutex_unlock(lock2);</pre>

Thread attributes

- ▶ Specific attributes of the thread can be set at create time using the `pthread_attr_t` object
 - ▶ Detached or joinable
 - ▶ Scheduling information
 - ▶ Stack details (size, address, overflow size)
- ▶ Need to use the `pthread_attr_init` and `pthread_attr_destroy` functions to managed the attribute object

Thread Safety

- ▶ Library calls often use a global data structure (seed of RNG)
- ▶ With multiple threads calling routines the data structure may not remain consistent
- ▶ The library needs to put in some synchronisation to prevent corruption
- ▶ This makes the library *thread safe*
- ▶ If using libraries and threads, make sure they are thread safe
- ▶ Otherwise you will have to serialize the function calls

Debugging Threads

- ▶ As always start debugging using `printf`
- ▶ Can do thread debugging in `gdb`
- ▶ `break function thread 2`
- ▶ This will stop the execution once thread 2 calls `function`
- ▶ Usually all threads will stop execution. However, on some platforms there is a mode where all other threads continue executing in the background.
- ▶ Be careful stepping through code — other threads may execute more than one instruction in the time it takes the thread of interest to step.