

What is OpenMP?

- ▶ OpenMP is an API for developing shared memory application
- ▶ Cross platform — supported by most vendors
- ▶ C, C++ and Fortran
- ▶ Based on thread model of programming
- ▶ Formalized in 1997

What is OpenMP?

- ▶ Compiler directives to generate parallel code
- ▶ Runtime library that implements the functionality
- ▶ Environment variables control parallel execution at runtime

Compiler Directives

- ▶ `#pragma omp directive [clause] [clause] ...`
- ▶ The directive is applied to the next statement — often a new block
- ▶ The directive can be split over more than one line if required by using a backslash

Hello world

```
#include <omp.h>
#include <stdio.h>

int main() {

#pragma omp parallel
{
    int id;
    id = omp_get_thread_num();
    printf ("%d: _Hello_world\n", id);
}

return 0;
}
```

► gcc -fopenmp hello_omp.c

A note on compilers

- ▶ Most modern compilers support OpenMP
- ▶ GNU — `gcc -fopenmp foo.c`
- ▶ Intel — `icc -openmp foo.c`
- ▶ Pathscale — `pathcc -openmp foo.c`
- ▶ Potential issue with gcc built using homebrew on Mac

Creating Threads

- ▶ By default OpenMP creates a number of threads equal to the number of processors in the system
- ▶ There are several ways to modify this
 - ▶ **Library call** — `omp_set_num_threads(3);`
 - ▶ **Directive** — `#pragma omp parallel num_threads(3)`
 - ▶ **Environment variable** — `% export OMP_NUM_THREADS=3`
- ▶ Can use `omp_get_num_procs()` to find out how many processors are available for you to use

Loop Parallelization

- ▶ Often the evaluation of statements in a loop are independant of each other
- ▶ The order of execution is not important
- ▶ These can be easily parallelized in OpenMP using the `for` pragma

```
#pragma omp parallel for  
for (i=0; i<N; i++)  
    a[i] = b[i]*c[i];
```

- ▶ The for loop conditions must be relatively normal
- ▶ No `break`, `return`, `exit`
- ▶ Each thread executes a section of the loop index range

Loop Parallelization

- ▶ What if we have a double loop?

```
#pragma omp parallel for
for(i=0;i<N;i++) {
    for(j=0;j<N;j++) {
        a[i][j] = b[i][j] * c[i][j];
    }
}
```

- ▶ Each thread will execute their own section of the `i` loop
- ▶ Since `j` is shared things will get very confused
- ▶ Could move declaration of `j` inside the `i` loop
- ▶ Could just do the `omp` on the inner loop — inefficient

Per Thread Variables

- ▶ Unless otherwise specified, all variables declared outside the block are shared
- ▶ One exception is the `for` loop index variable
- ▶ Variables declared inside the block are unique for that thread
- ▶ Make variable private to the thread by modifying the construct
- ▶ `#pragma omp parallel private(var)`
- ▶ This makes a new copy of the variable `var` for each thread
- ▶ The value of the new `var` is undefined

Loop Parallelization

- Now our double loop works fine

```
#pragma omp parallel for private(j)
for(i=0;i<N;i++) {
    for(j=0;j<N;j++) {
        a[i][j] = b[i][j] * c[i][j];
    }
}
```

Per Thread Variables

- ▶ `#pragma omp parallel firstprivate(var)`
- ▶ This will create a new variable for each thread but initialize to the value of the already declared variable `var`
- ▶ Using `private` and `firstprivate`, at the end of the block the old value of `var` is unchanged

```
var = 101;
#pragma omp parallel for firstprivate(var)
for (i=0; i<N; i++) {
    printf("var_=_%d\n", var);
    var = i;
}
printf("var_=_%d\n", var); /* This will print 101 */
```

Per Thread Variables

- ▶ Sometimes we want to be able to use the value of variables after the loop has completed
- ▶ Adding a `lastprivate` clause to the construct will set the value of the variable after the parallel region to the value it would have, if the loop had been executed serially

```
var = 4;  
printf("%d\n", var);
```

```
#pragma omp parallel for lastprivate(var)  
for(i=0;i<N;i++) {  
    var = i;  
}
```

```
printf("%d\n", var); /* Will print value N-1 */
```

Accessing Shared Variables

- ▶ As with pthreads, need a way to control updates to shared variables
- ▶ Can introduce race conditions that give the wrong answer

```
max = INT_MIN;  
#pragma omp parallel for  
for (i=0; i<N; i++) {  
    if (a[i] > max)  
        max = a[i];  
}
```

- ▶ Without OpenMP, this loop finds the maximum value in an array
- ▶ With multiple threads, a max value may be skipped due to order of read/write instructions in different threads

Critical Sections

- ▶ Can fix this using a `critical` directive
- ▶ Only one thread can be executing in a `critical` section at a time
- ▶ This version, while correct is terribly inefficient
- ▶ We have pretty much serialized the code again

```
max = INT_MIN;
#pragma omp parallel for
for(i=0;i<N;i++) {
    #pragma omp critical
    {
        if(a[i] > max)
            max = a[i];
    }
}
```

Critical Sections

- ▶ We can't just move the `critical` directive to after the test
- ▶ The value of `max` increases monotonically
- ▶ Most loop trips will not update `max`

```
max = INT_MIN;
#pragma omp parallel for
for(i=0;i<N;i++) {
    if(a[i] > max) {
        #pragma omp critical
        {
            if(a[i] > max)
                max = a[i];
        }
    }
}
```

Reduction Clauses

- ▶ Often we want to carry out a reduction on a variable in a loop
- ▶ OpenMP allows us to add this to the directive
- ▶ Sum, product, bitwise and logical operations

critical	reduction
<pre>total = 0; #pragma omp parallel for for(i=0;i<N;i++) { a[i] = some_func(); #pragma omp critical total += a[i]; }</pre>	<pre>total = 0; #pragma omp parallel for reduction(+:total) for(i=0;i<N;i++) { a[i] = some_func(); total += a[i]; }</pre>

Conditional Loop Parallelization

- ▶ There is an overhead in creating and joining threads
- ▶ If the trip count of a loop is too small it may just be quicker to execute in serial
- ▶ We can add a condition to the directive to decide at runtime whether or not to parallelize the loop
- ▶ Need to do some experiments to determine what the threshold should be

```
total = 0;
#pragma omp parallel for reduction(+:total) if(N > 200)
for(i=0;i<N;i++) {
    a[i] = some_func();
    total += a[i];
}
```

Work Scheduling

- ▶ With the default `for` division of work between the threads we can have imbalances
- ▶ For example - initialising an upper triangular matrix
- ▶ Thread 0 will have $N + N - 1 + N - 2 \dots$ calls to `func` whereas the last thread will have $\dots 3 + 2 + 1$ calls to `func`

```
#pragma omp parallel for private(j)
for(i=0;i<N;i++)
    for(j=i;j<N;j++)
        a[i][j] = func(i, j);
```

Work Scheduling

- ▶ We can have three types of schedule in OpenMP
 - ▶ static — all iterations are assigned to a thread before the start execution
 - ▶ dynamic — threads execute some iterations and are then assigned more until all work is complete
 - ▶ guided — like dynamic but the size of work packages changes over time heuristically
- ▶ We also define a `chunk` to be a range of contiguous iterations of a loop
- ▶ Now we can add the clause `schedule(type, chunk)` to the OpenMP directive. The `chunk` is optional

Work Scheduling

- ▶ `schedule(static)` — every thread gets about n/t iterations
- ▶ `schedule(static, N)` — every thread gets given an interleaved allocation of sets of N iterations
- ▶ `schedule(dynamic)` — threads are given iterations one at a time
- ▶ `schedule(dynamic, N)` — threads are given iterations N at a time
- ▶ `schedule(guided)` — the system uses a heuristic function to allocate threads down to chunks of size one
- ▶ `schedule(guided, N)` — as guided but to minimum chunk size of N
- ▶ `schedule(runtime)` — looks at the environment variable `OMP_SCHEDULE` to determine what to use

Backward Compatability

- ▶ Even though most compilers now support OpenMP you may for some reason want to compile and run your code without OpenMP
- ▶ Recall there are two ways your code has been modified
 - ▶ Directives
 - ▶ Library Calls
- ▶ If the compiler reaches a directive it doesn't understand it silently ignores it
- ▶ However you will need to work around any library calls

```
#ifdef _OPENMP
    nthreads = omp_get_num_procs();
#else
    nthreads = 1;
#endif
```

Programming Models

- ▶ Pure MPI — one MPI task on each core
- ▶ Pure OpenMP — limited to 1 machine or use distributed shared memory
- ▶ Hybrid — MPI for internode, OpenMP within the node
 - ▶ No overlap of comms — MPI outside OpenMP regions
 - ▶ Overlap of comms — MPI within OpenMP regions
- ▶ Code may not be portable. Different MPI implementations support threads in different ways

Why mix MPI and OpenMP

- ▶ Modern compute nodes have many cores — 8 to 24 being common
- ▶ Still tend to only have a single network connection
- ▶ With 24 MPI tasks per node there will be lots of contention for the available network resources
- ▶ Might be better to reduce the amount of comms outside the node — 2 MPI tasks with 12 threads or 4 MPI tasks with 6 threads
- ▶ Avoids unnecessary overhead doing MPI within the node

Support for Hybrid programming

- ▶ Your MPI implementation may have limited or no support for hybrid programming
- ▶ There are four levels of thread support
 - ▶ `MPI_THREAD_SINGLE` — only one thread
 - ▶ `MPI_THREAD_FUNNELED` — only the main thread can issue MPI calls
 - ▶ `MPI_THREAD_SERIALIZED` — all threads can issue MPI calls but only one thread at a time
 - ▶ `MPI_THREAD_MULTIPLE` — all threads can issue MPI calls at any time
- ▶ Each is represented by an integer and they are strictly ordered. The higher the value the more support there is
- ▶ Replace `MPI_Init` with `MPI_Init_thread` to see what your MPI supports

Determining your level of support

```
requested = MPI_THREAD_MULTIPLE;  
MPI_Init_thread(&argc, &argv, requested, &provided);  
  
switch(provided) {  
    case MPI_THREAD_SINGLE:  
        printf("Support for MPI_THREAD_SINGLE\n");  
        break;  
    case MPI_THREAD_FUNNELED:  
        printf("Support for MPI_THREAD_FUNNELED\n");  
        break;  
    case MPI_THREAD_SERIALIZED:  
        printf("Support for MPI_THREAD_SERIALIZED\n");  
        break;  
    case MPI_THREAD_MULTIPLE:  
        printf("Support for MPI_THREAD_MULTIPLE\n");  
        break;  
}
```

Mixing MPI and OpenMP

- ▶ In order to mix MPI and OpenMP we need at least `MPI_THREAD_FUNNELED` support
- ▶ We build our MPI application and then use OpenMP directives to parallelize the compute parts
- ▶ MPI calls are only used in the main MPI thread
- ▶ We can use the `MPI_Is_thread_main(int *flag)` to check if we are the thread that is allowed to make MPI calls

Other thread libraries

- ▶ We can use other thread libraries such as `pthread`s instead of OpenMP to spawn threads
- ▶ More flexible but introduces the usual problem of controlling access to shared data
- ▶ Best if MPI configured to provide `MPI_THREAD_MULTIPLE` support
- ▶ Note MPI Send/Recv can address only a process, not a thread within the process
- ▶ Can use tags as an additional identifier
- ▶ Be very careful when mixing MPI, threads and non-blocking calls
- ▶ Data will probably not end up where you expect it to!!